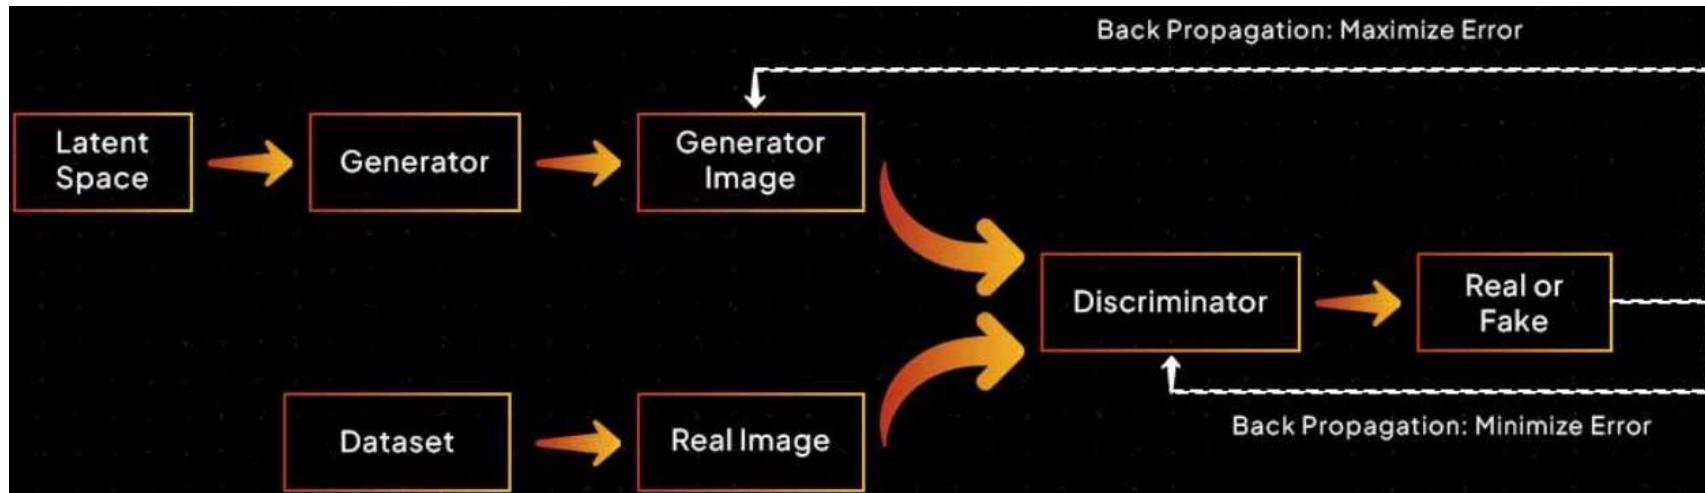


CIFAR10 USING GAN

Presented By Akshit Jain

Implement generator and the discriminator architectures by introducing convolution and up/transposed convolution layers. Use one of the classes from the CIFAR-10 dataset to generate images. Show the effect of controlling the noise vector.

GAN Architecture



CIFAR-10 Classification

- **Dataset:** CIFAR-10 has 60,000 32x32 RGB images
- **Output Class:** 10 Classes
- **Frequency:** Each class has exactly 5,000 rows.

Training Setup

- **Optimizer:** Adam
- **Epochs:** 5000
- **Batch Size:** 64
- **Loss:** BinaryCrossentropy
- **Target Class:** 7 (Horse) – Only used 1 class

Generator

- **Input :** Random noise vector (latent_dim = 100)
- **Dense Expansion :** Transforms noise into 8 x 8 x 256 feature.
- **Reshape & Activation:** Reshapes to 8 x 8 x 256. Used LeakyReLU
- **BatchNormalization:** Yes
- **UpSampling (8 x 8 → 8 x 8) :** Uses Conv2DTranspose to refine features
- **Resolution Increase (8 x 8 → 16 x 16):** Another Conv2DTranspose doubles the size.
- **Final Output (16 x 16 → 32 x 32 x 3):** Last Conv2DTranspose creates 32 x 32 RGB image with Tanh activation

Generator

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 100)	0
dense (Dense)	(None, 16384)	1,638,400
batch_normalization (BatchNormalization)	(None, 16384)	65,536
leaky_re_lu (LeakyReLU)	(None, 16384)	0
reshape (Reshape)	(None, 8, 8, 256)	0
conv2d_transpose (Conv2DTranspose)	(None, 8, 8, 128)	819,200
batch_normalization_1 (BatchNormalization)	(None, 8, 8, 128)	512
leaky_re_lu_1 (LeakyReLU)	(None, 8, 8, 128)	0
conv2d_transpose_1 (Conv2DTranspose)	(None, 16, 16, 64)	284,800
batch_normalization_2 (BatchNormalization)	(None, 16, 16, 64)	256
leaky_re_lu_2 (LeakyReLU)	(None, 16, 16, 64)	0
conv2d_transpose_2 (Conv2DTranspose)	(None, 32, 32, 3)	4,800

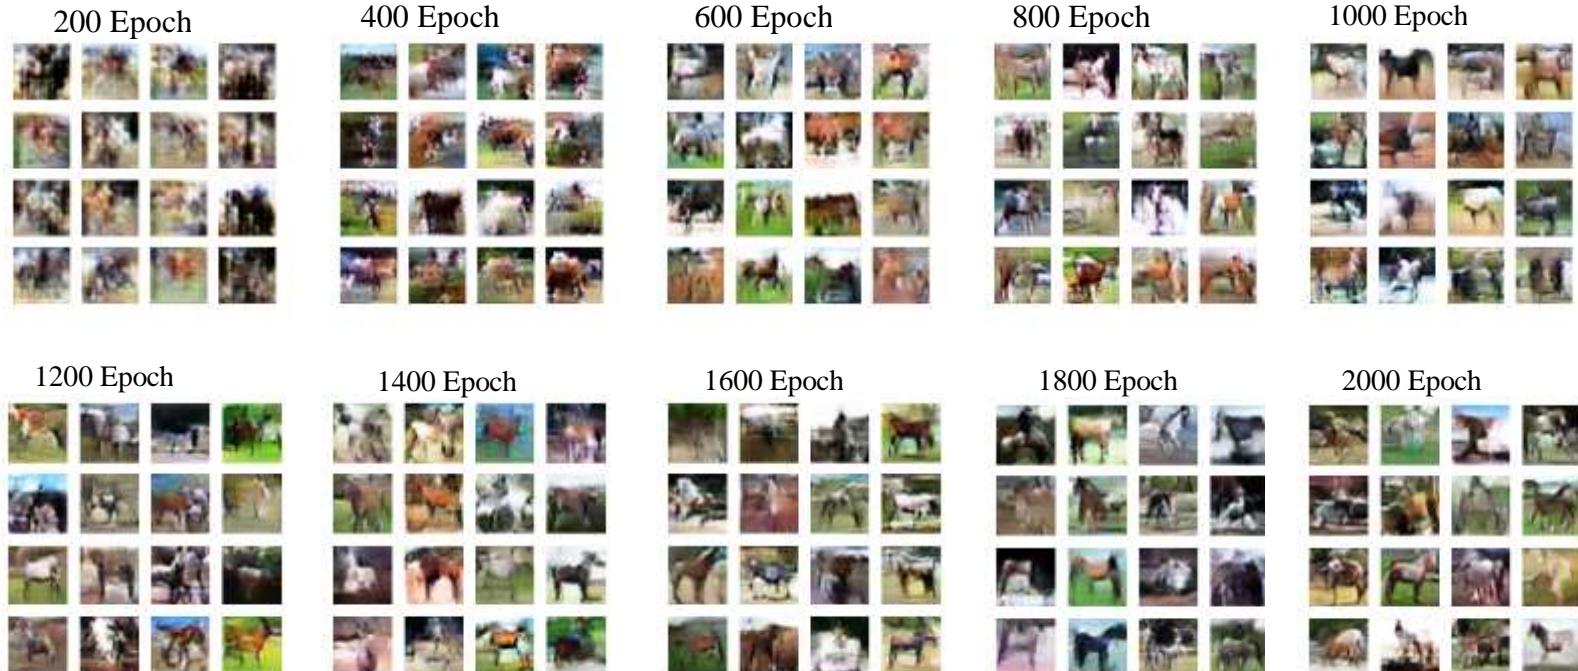
Discriminator

- **Input:** Takes a $32 \times 32 \times 3$ image as input.
- **First Convolution ($32 \times 32 \rightarrow 16 \times 16$):** Extracts features with a 5×5 Conv2D layer, using stride 2 for downsampling.
- **Activation & Regularization:** LeakyReLU introduces non-linearity, followed by Dropout (0.3) to prevent overfitting.
- **Second Convolution ($16 \times 16 \rightarrow 8 \times 8$):** Another 5×5 Conv2D further downsamples and extracts deeper features.
- **Flatten & Dense Layer:** Flattens the feature maps into a 1D vector and passes it to a single Dense layer.
- **Output:** Outputs a single value (logit) indicating whether the image is real or fake.

Discriminator

Layer (type)	Output Shape	Param #
input_layer_1 (InputLayer)	(None, 32, 32, 3)	0
conv2d (Conv2D)	(None, 16, 16, 64)	4,864
leaky_re_lu_3 (LeakyReLU)	(None, 16, 16, 64)	0
dropout (Dropout)	(None, 16, 16, 64)	0
conv2d_1 (Conv2D)	(None, 8, 8, 128)	204,928
leaky_re_lu_4 (LeakyReLU)	(None, 8, 8, 128)	0
dropout_1 (Dropout)	(None, 8, 8, 128)	0
flatten (Flatten)	(None, 8192)	0
dense_1 (Dense)	(None, 1)	8,193

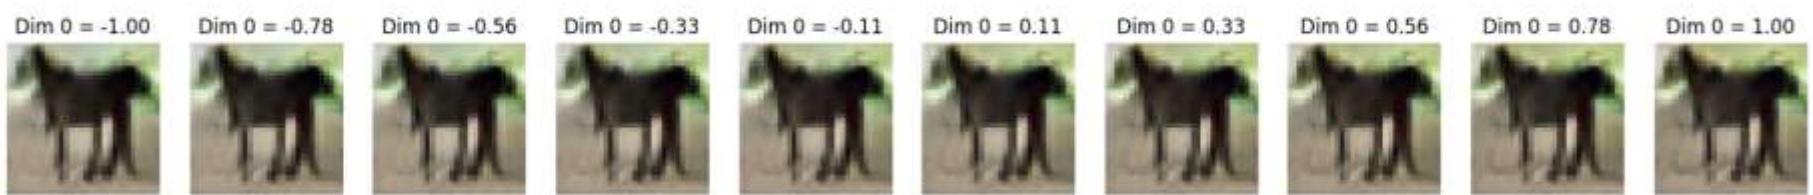
Generated images from class 7 (Horse)



Results

Effect of Noise vector

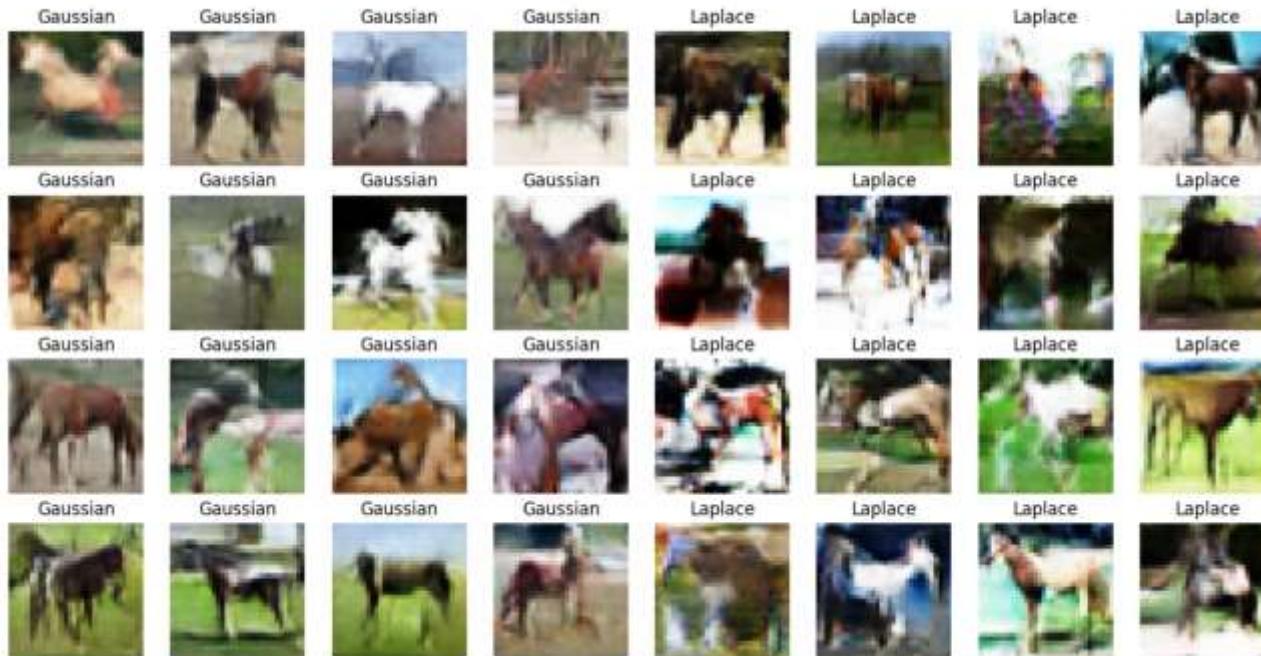
Effect of Varying Latent Dimension 0



Results

Comparison of Gaussian and Laplace Noise in Latent Space

Comparison of Gaussian and Laplace Noise in Latent Space



Thank You